



# Buffering Techniques

---

*Greg Stitt*  
*ECE Department*  
*University of Florida*



# Buffers

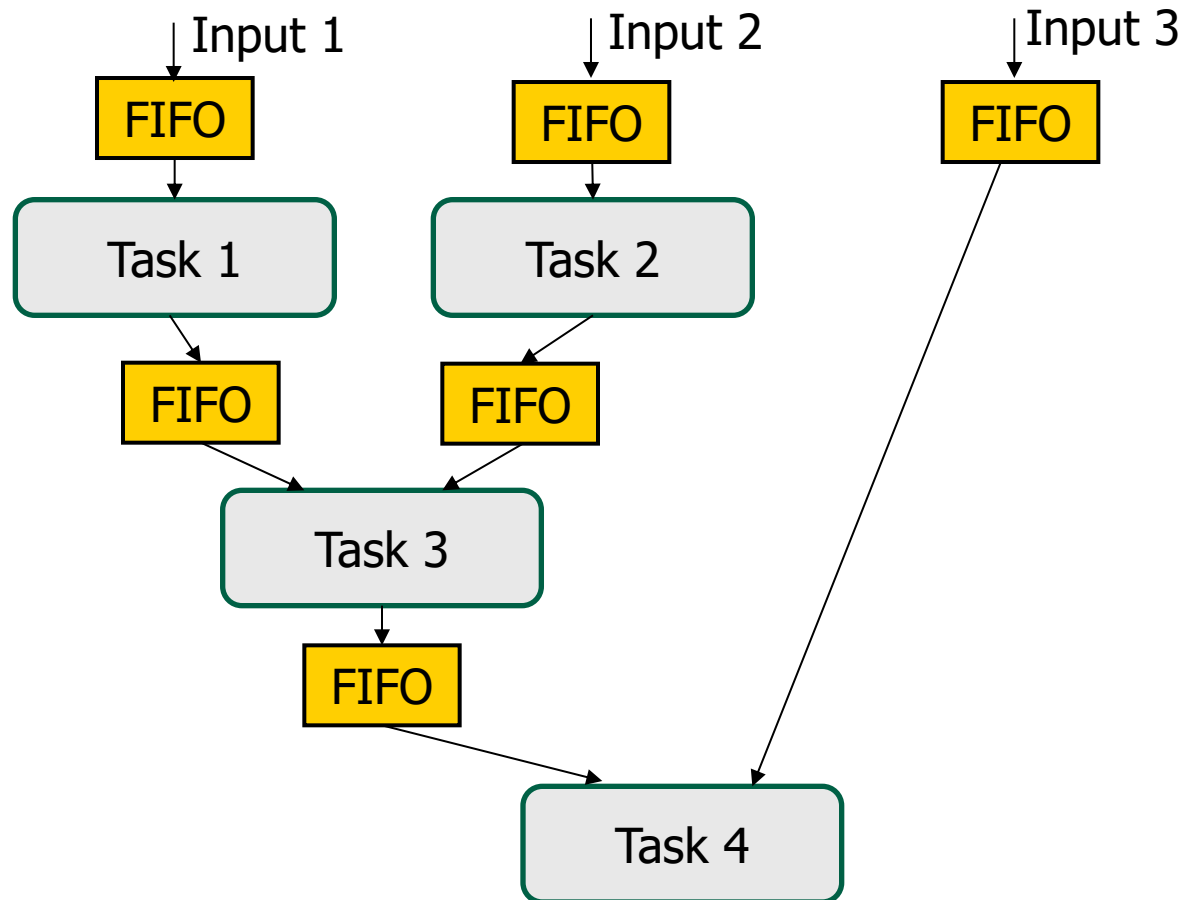
---

- Purposes

- Flow control for tasks with different production and consumption rates
- Metastability issues
  - Memory clock likely different from circuit clock
  - Buffer stores data at one speed, circuit reads data at another
- Stores “windows” of data, delivers to datapath
  - Window is set of inputs needed each cycle by pipelined circuit
  - Generally, more efficient than datapath requesting needed data
    - i.e. Push data into datapath as opposed to pulling data from memory
- Conversion between memory and datapath widths
  - E.g. Bus is 64-bit, but datapath requires 128-bits every cycle
    - Input to buffer is 64-bit, output from buffer is 128 bits
    - Buffer doesn't say it has data until receiving pairs of inputs

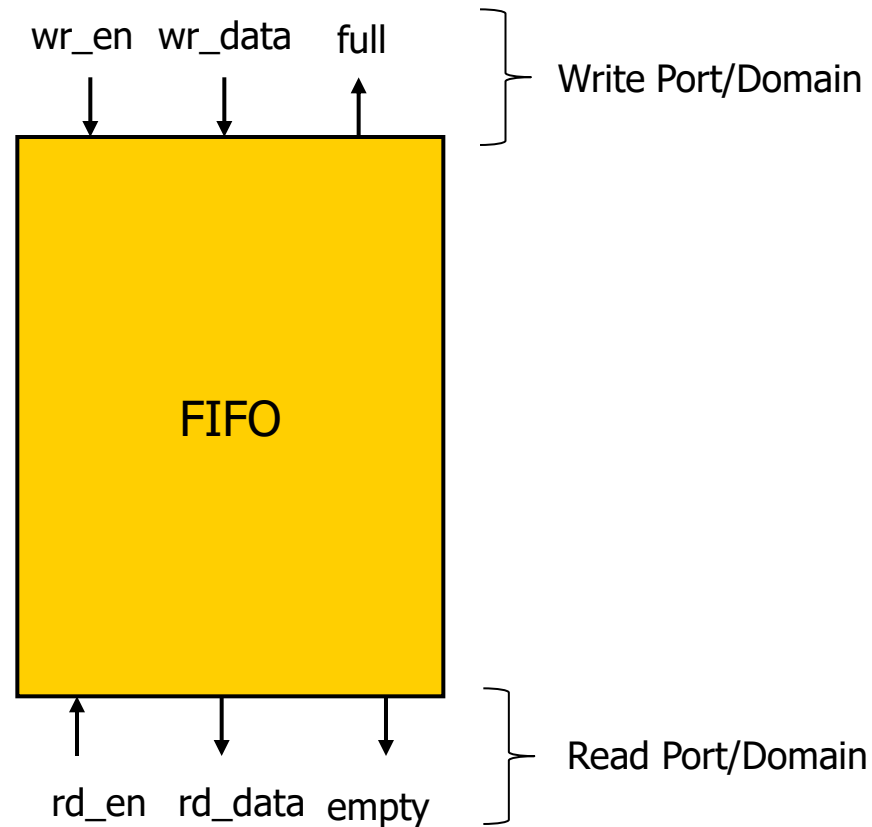
# FIFOs

- Most common buffer: First-in First-out (FIFO)
  - Used for synchronizing execution of parallel tasks
    - i.e., “flow control”



# FIFOs

- Beautifully simple control
  - Producer writes if not full (full provides “back pressure”)
  - Consumer reads if not empty





# FIFO Sizing

---

- Tasks often produce/consume data at different rates
  - Tasks can also just be streams of data
  - e.g., Ethernet data, AXI bus data, memory data
- We often need to “size” the FIFO accordingly to handle different rates without back pressure
  - $F_p$  = max message rate of producer (msg/s)
  - $F_c$  = max message rate of consumer (msg/s)
  - We use abstract message to generalize the problem



# FIFO Sizing

---

- Three common situations:
  - 1)  $F_p \leq F_c$ , assuming consumer is always ready
    - FIFO only needed if tasks on different clock domains
      - FIFO depth of 1 (or 0)
    - Why traditional pipelines use registers instead of FIFOs
      - The execution between pipeline operations (i.e. tasks) is synchronized implicitly every clock cycle, no need for a FIFO
  - 2)  $F_p > F_c$ , assuming a constant producer stream
    - FIFO would have to be infinitely large to avoid filling up
    - If this is unavoidable, producer must monitor the full flag
      - i.e., “back pressure” to slow down a producer



# FIFO Sizing

---

- 3)  $F_p > F_c$ , but producer sends data in bursts
  - E.g., Ethernet capable of 10/40/100 Gb/s, but data not available every cycle
  - We'll use  $B$  to represent the maximum number of messages that a producer might send consecutively (i.e., a burst)
  - We'll also assume that after a burst, the FIFO has time to empty
- How to determine FIFO depth that avoids back pressure?
  - $F_p - F_c =$  rate that FIFO fills up
    - e.g., after 1 second, FIFO would contain  $F_p - F_c$  messages
  - Next, we need to combine the "fill rate" with the duration of the burst  $B$ 
    - $B = \#$  of consecutive messages
    - $F_p = \text{msg} / \text{s}$
    - Burst duration =  $B / F_p$



# FIFO Sizing

---

- FIFO depth = fill\_rate \* burst\_duration
$$= (F_p - F_c) * B / F_p$$
$$= B * (1 - F_c / F_p)$$
- E.g.:
  - $F_p = 1\text{M msgs / s}$
  - $F_c = 250\text{k msgs / s}$
  - $B = 1000\text{ msgs}$
  - FIFO Depth = 750
- Remember: you could also use a smaller depth with back pressure





# FIFO Sizing

---

- Similar problem


- If time between bursts is sufficient for FIFO to drain, this means producer data rate and consumer data rate are equal within *some* window of time
- $D_p$  = Producer data rate (msgs / window)
- $D_c$  = Consumer data rate (msgs / window)
- $D_p = D_c$
- But, don't know when data produced/consumed within window

- What is the minimum FIFO depth to avoid back pressure?

- Sliding window of time
  - Can't produce more than  $D_p$  messages within *any* window
- Worst case situation:
  - All  $D_p$  messages arrive at beginning of window
  - All  $D_c$  messages are consumed at end of window


# FIFO Sizing

- Consider several situations:
- 1) Window is much longer than time to produce/consume all messages



# FIFO Sizing

- 2) Window small enough to create overlap between Dp and Dc




$$\text{Min Depth} = \text{Window} - Dc$$

- Two situations can be combined:
  - **FIFO depth = min(window - Dc, Dp)**
- Examples for Dp = Dc = 10 messages:
  - Window = 100 cycles => FIFO depth = 10
  - Window = 15 cycles => FIFO depth = 5
  - Window = 10 cycles => FIFO depth = 0 (not needed unless different clocks)

# FIFOs

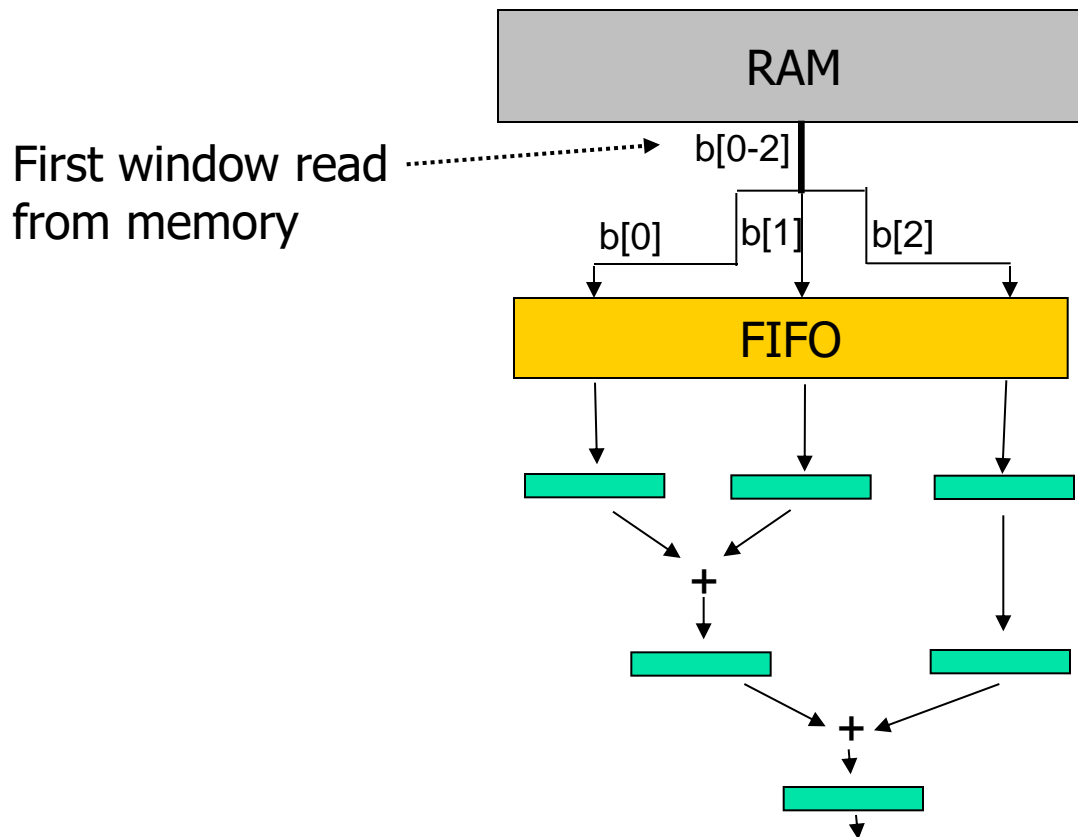
- FIFOs often used between RAM and pipeline
  - Outputs data in order read from memory



```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

# FIFOs


- FIFOs often used between RAM and pipeline
  - Outputs data in order read from memory



```
for (i=0; i < 100; I++)  
    a[i] = b[i] + b[i+1] + b[i+2];
```

# FIFOs

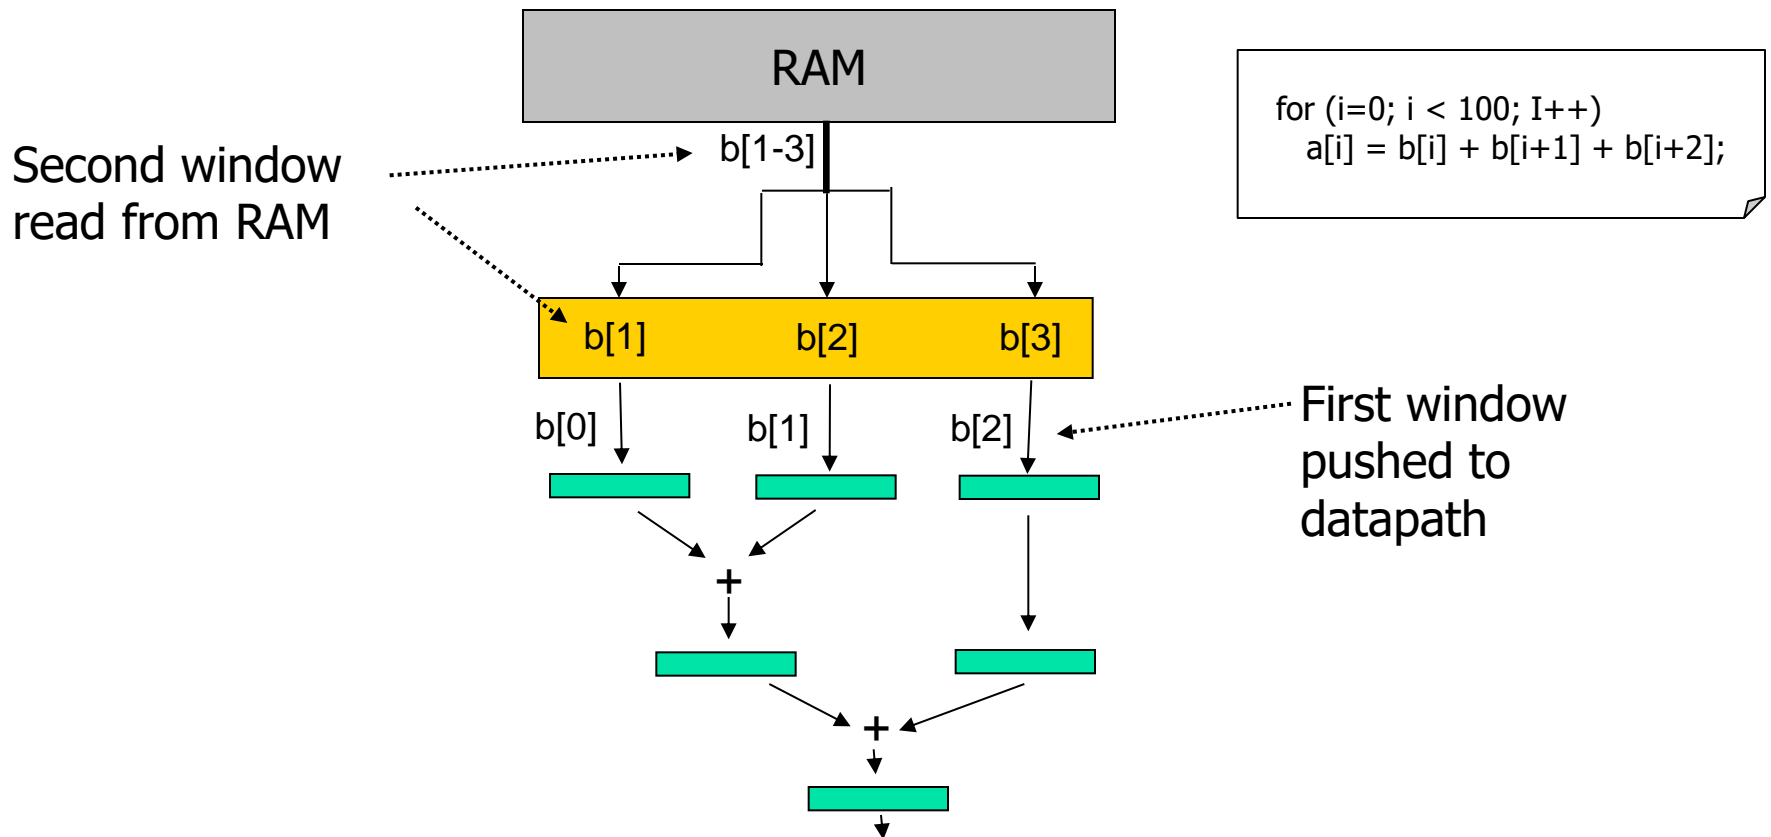
- FIFOs often used between RAM and pipeline
  - Outputs data in order read from memory



```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

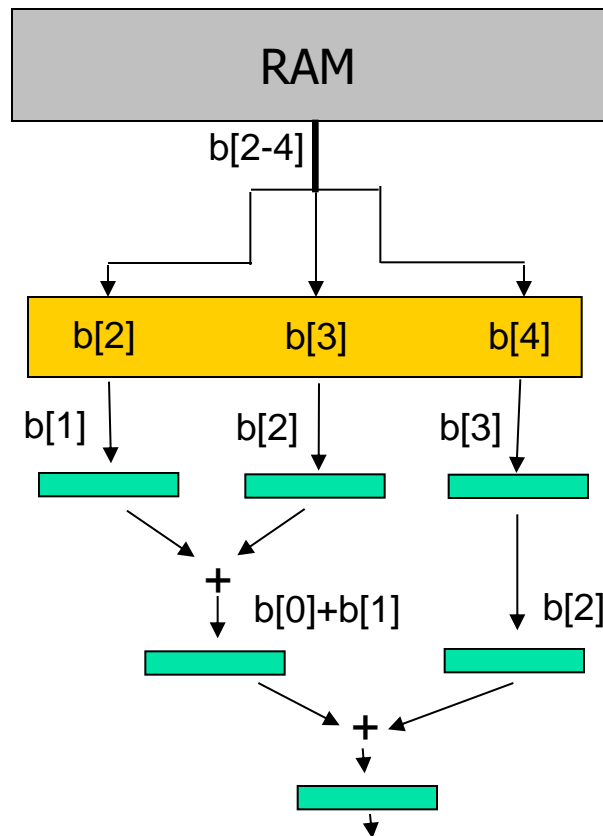
# FIFOs

- FIFOs often used between RAM and pipeline
  - Outputs data in order read from memory



# FIFOs

- FIFOs often used between RAM and pipeline
  - Outputs data in order read from memory



```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```





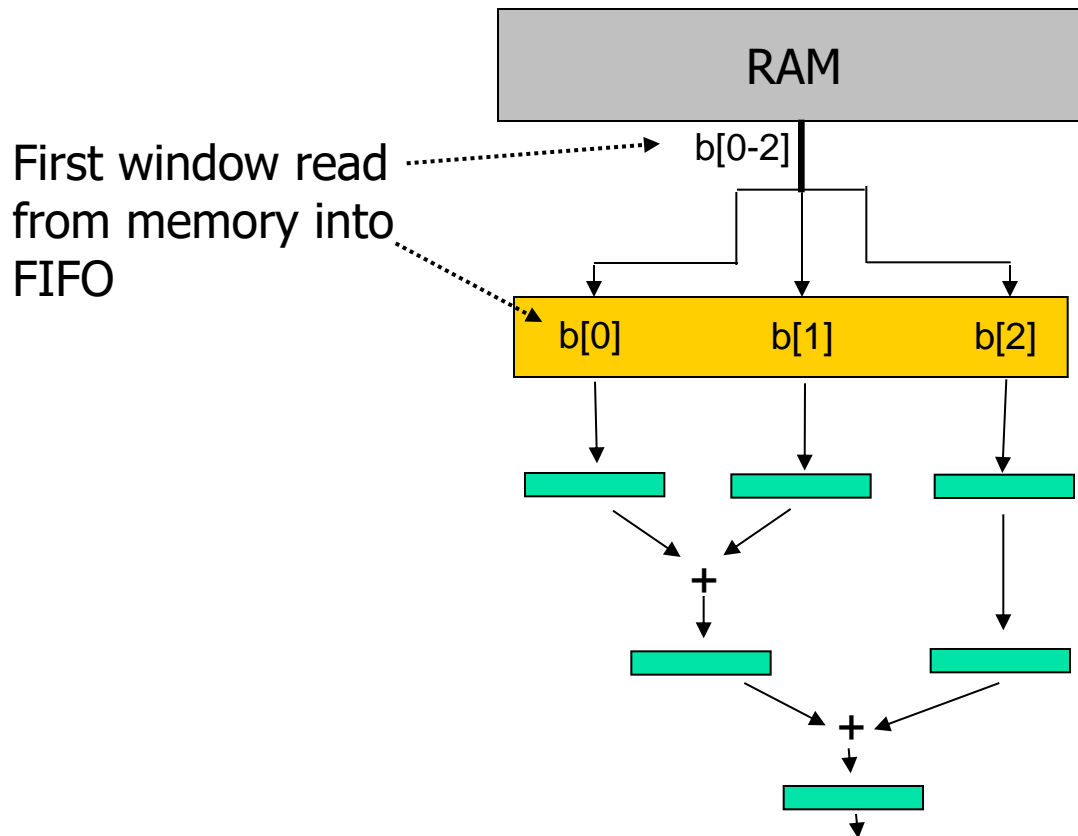
# FIFOs

---

- Timing issues
  - Memory bandwidth too small
    - Circuit stalls or wastes cycles while waiting for data
  - Memory bandwidth larger than data consumption rate of circuit
    - May happen if area exhausted

# FIFOs

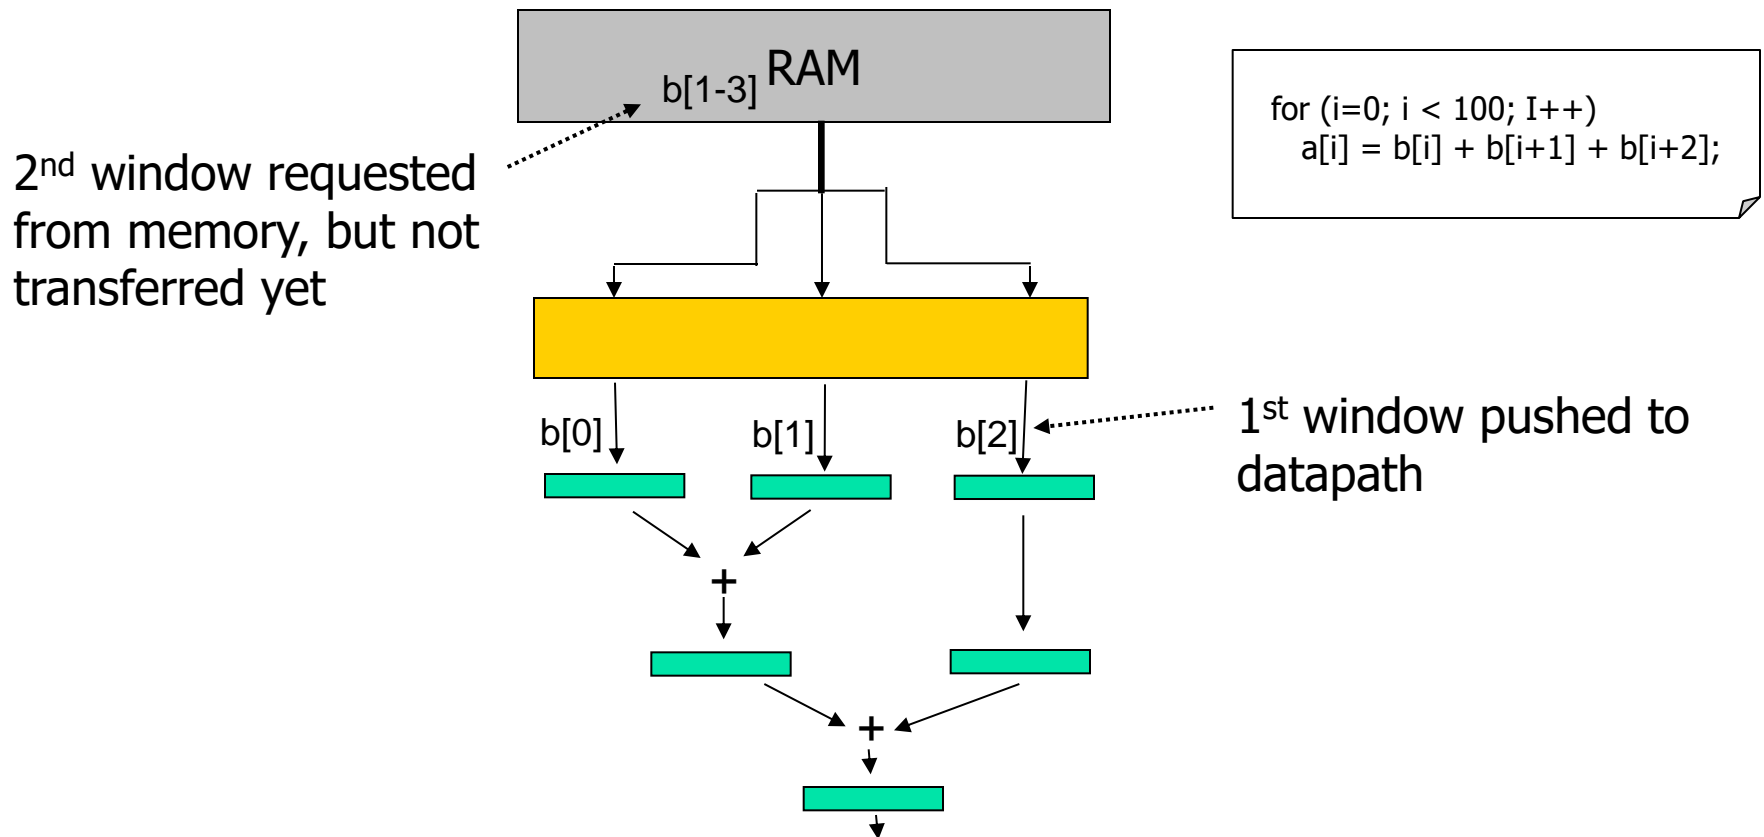
- Memory bandwidth too small



```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

# FIFOs


- Memory bandwidth too small



# FIFOs

- Memory bandwidth too small

2<sup>nd</sup> window requested from memory, but not transferred yet




```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

No data ready (wasted cycles)

Alternatively, could have prevented 1<sup>st</sup> window from proceeding (stall cycles) - *necessary if feedback in pipeline*

# FIFOs

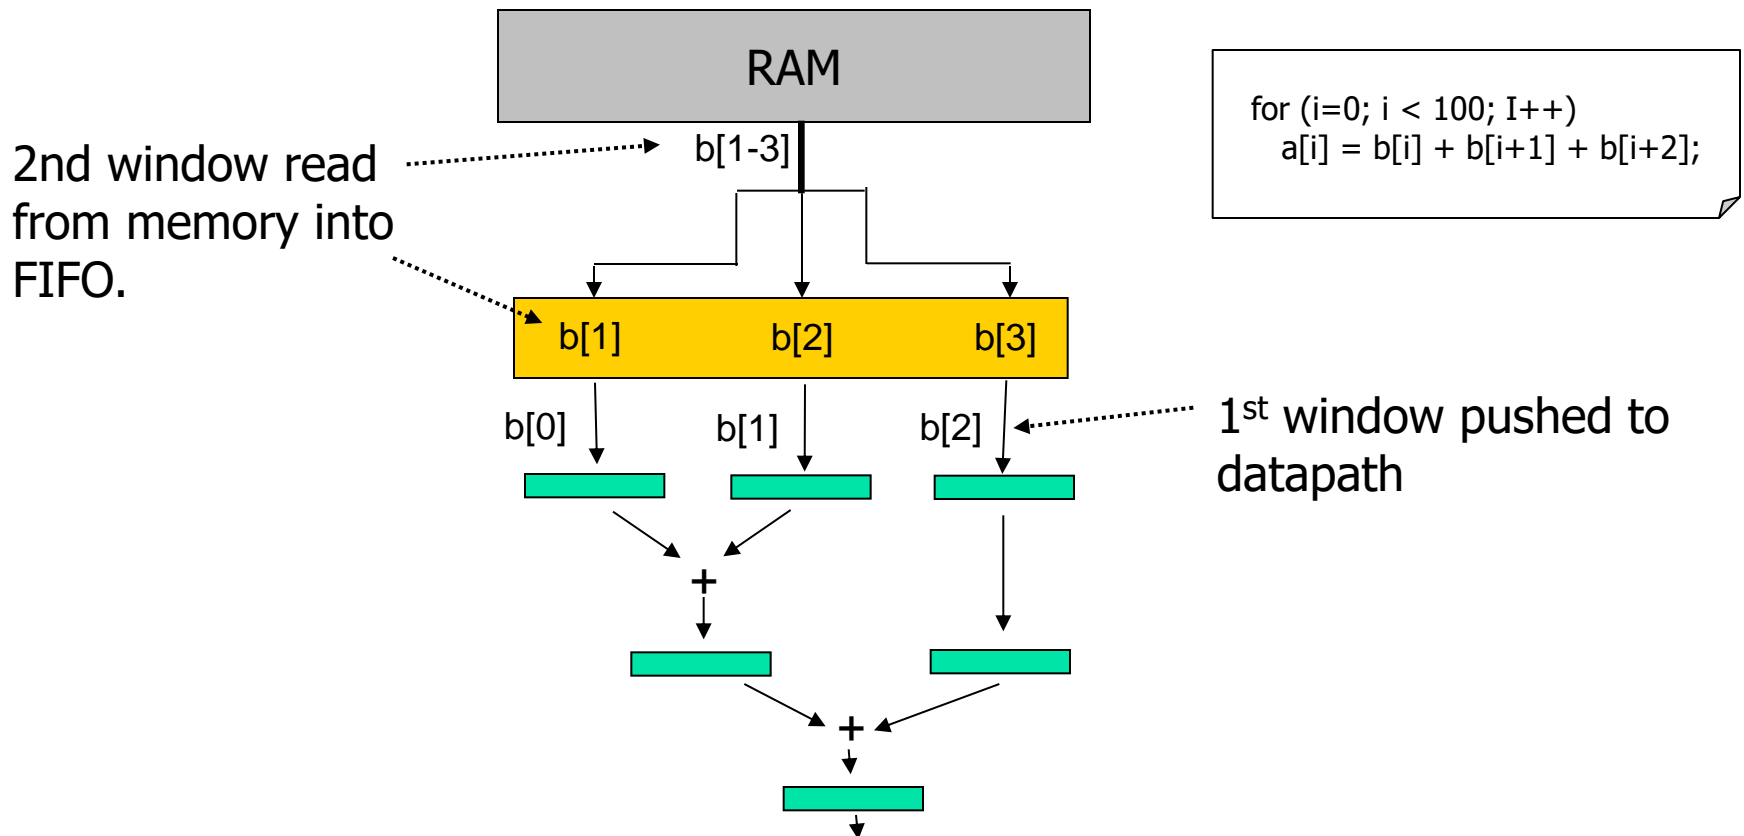
- Memory bandwidth larger than data consumption rate



```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```


# FIFOs

- Memory bandwidth larger than data consumption rate



# FIFOs

- Memory bandwidth larger than data consumption rate



```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```


..... Data arrives faster than  
circuit can process it –  
FIFOs begin to fill

If FIFO full, memory  
reads stop until not full

# Improvements

- Do we need to fetch entire window from memory for each iteration?
  - Only if windows of consecutive iterations are mutually exclusive
- Commonly, consecutive iterations have overlapping windows
  - Overlap represents “reused” data
  - Ideally, would be fetched from memory just once
- Smart Buffer [*Guo, Buyukkurt, Najjar LCTES 2004*]
  - Part of the ROCCC compiler
  - Analyzes memory access patterns
  - Detects “sliding windows”, reused data
  - Prevents multiple accesses to same data


```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```





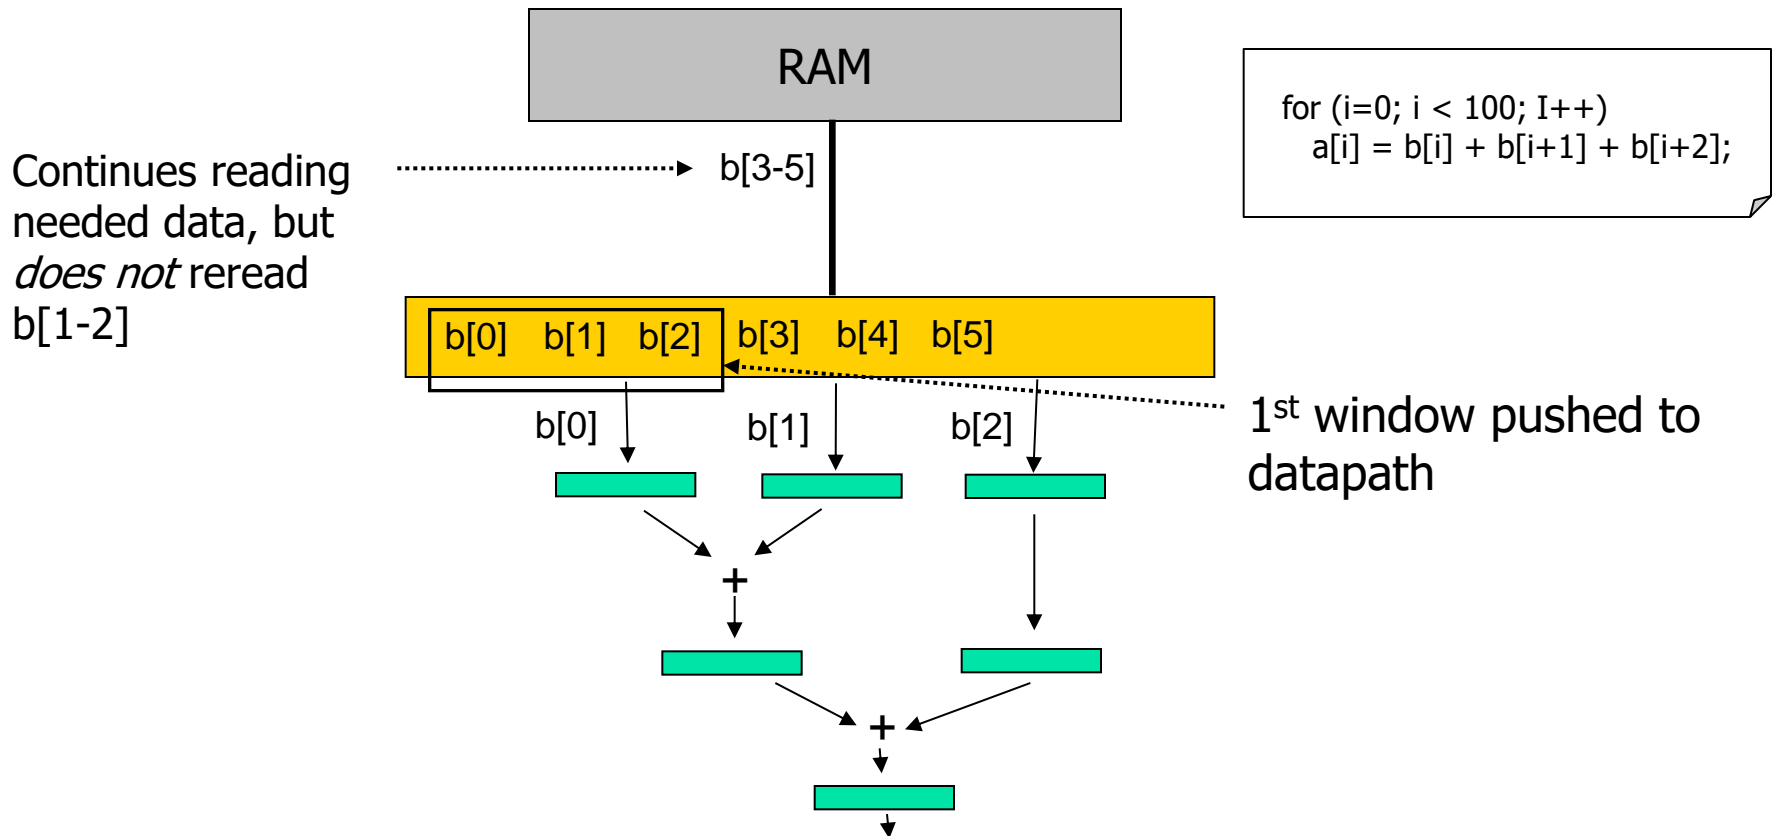
# Smart Buffer

1st window read  
from memory into  
smart buffer

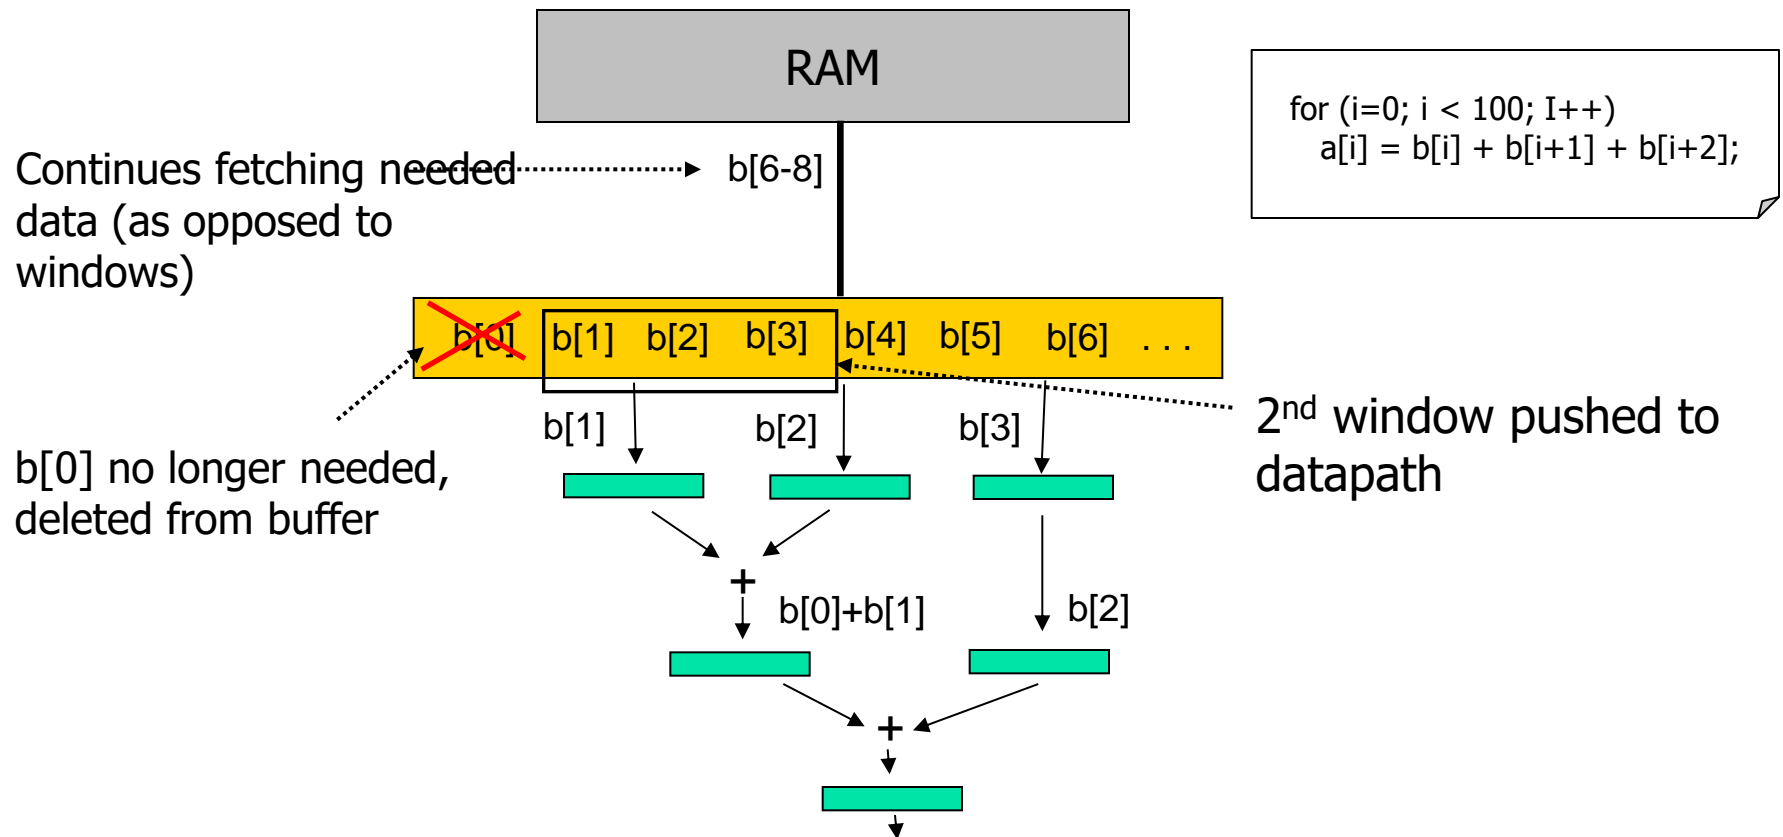


```
for (i=0; i < 100; I++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

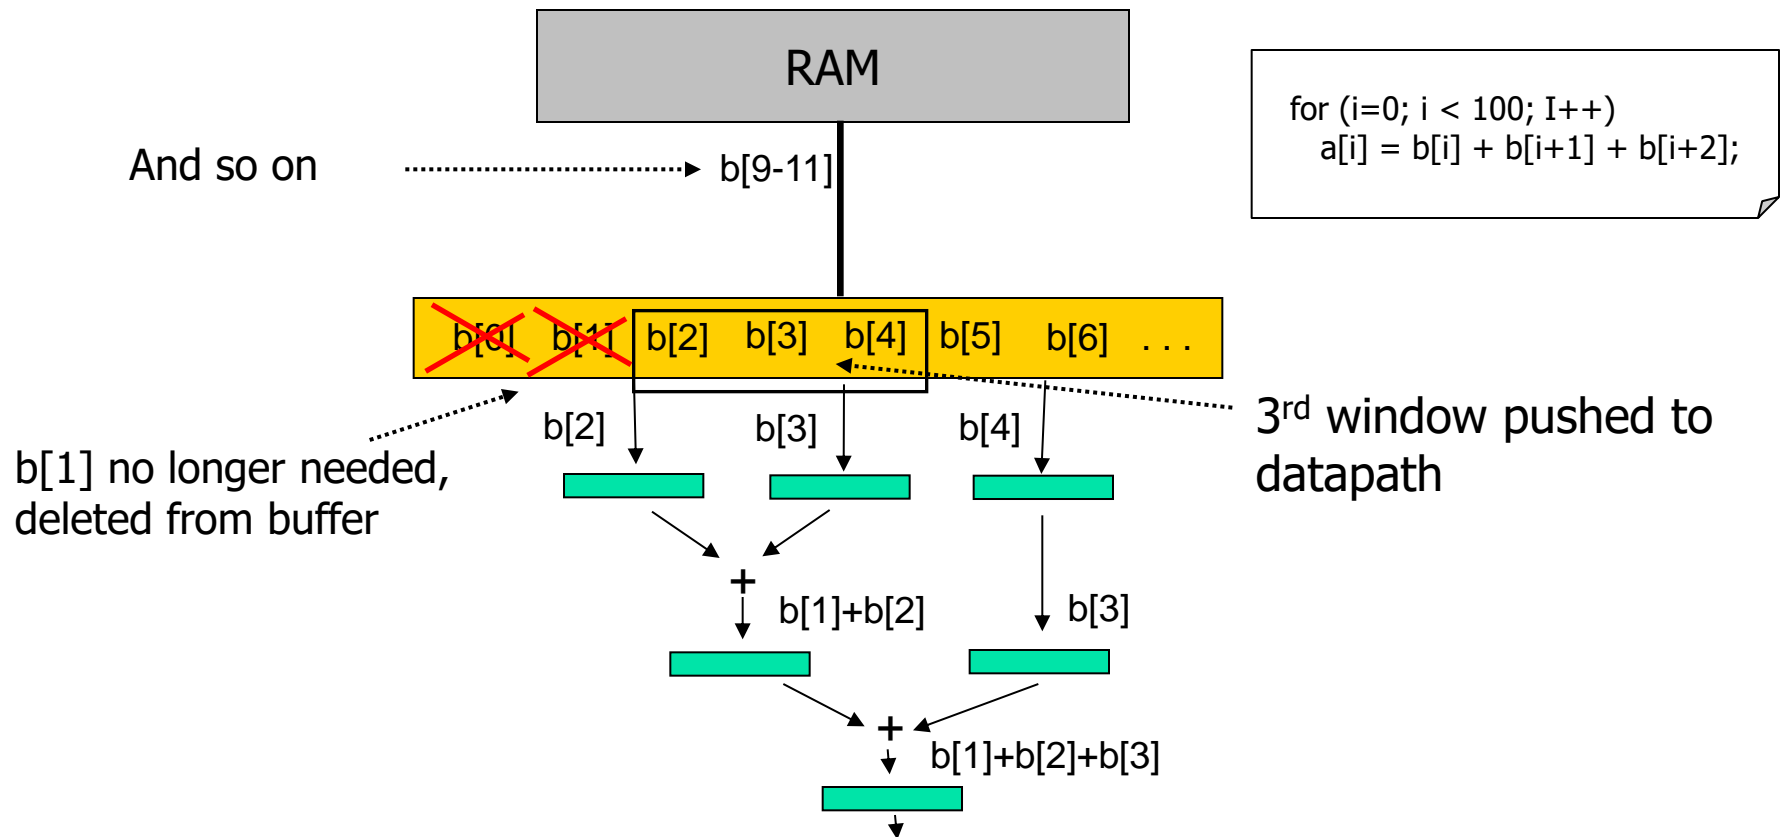
# Smart Buffer



# Smart Buffer



# Smart Buffer






# Comparison with FIFO

- FIFO - *fetches a window for each iteration*
  - Reads 3 elements every iteration
    - $100 \text{ iterations} * 3 \text{ accesses/iteration} = 300 \text{ memory accesses}$
- Smart Buffer - *fetches as much data as possible each cycle, buffer assembles data into windows*
  - Ideally, reads each element once
    - # accesses = array size
    - Circuit performance is equal to latency plus time to read array from memory
      - **No matter how much computation, execution time is approximately equal to time to stream in data!**
        - Note: Only true for streaming examples
  - 102 memory accesses
    - Essentially improves memory bandwidth by 3x
    - How does this help?
      - **Smart buffers enable more unrolling**

# Unrolling with Smart Buffers

- Assume bandwidth = 128 bits
  - We can read  $128/32 = 4$  array elements each cycle
  - First access doesn't save time
    - No data in buffer
    - Same as FIFO, can unroll once

```
long b[102];  
for (i=0; i < 100; I++)  
    a[i] = b[i] + b[i+1] + b[i+2];
```




First memory access allows only 2 parallel iterations

# Unrolling with Smart Buffers

- After first window is in buffer
  - Smart Buffer Bandwidth = Memory Bandwidth + Reused Data
    - 4 elements + 2 reused elements ( $b[2], b[3]$ )
    - Essentially, provides bandwidth of 6 elements per cycle
  - Can perform 4 iterations in parallel
    - **2x speedup compared to FIFO**

```
long b[102];  
for (i=0; i < 100; i++)  
    a[i] = b[i] + b[i+1] + b[i+2];
```



However, every subsequent access enables 4 parallel iterations ( $b[2]$  and  $b[3]$  already in buffer)



# Datapath Design w/ Smart Buffers

---

- Datapath based on unrolling enabled by smart buffer bandwidth (not memory bandwidth)
  - Don't be confused by first memory access
- Smart buffer waits until initial windows in buffer before passing any data to datapath
  - Adds a little latency
    - But, avoids 1<sup>st</sup> iteration requiring different control due to less unrolling





# Another Example

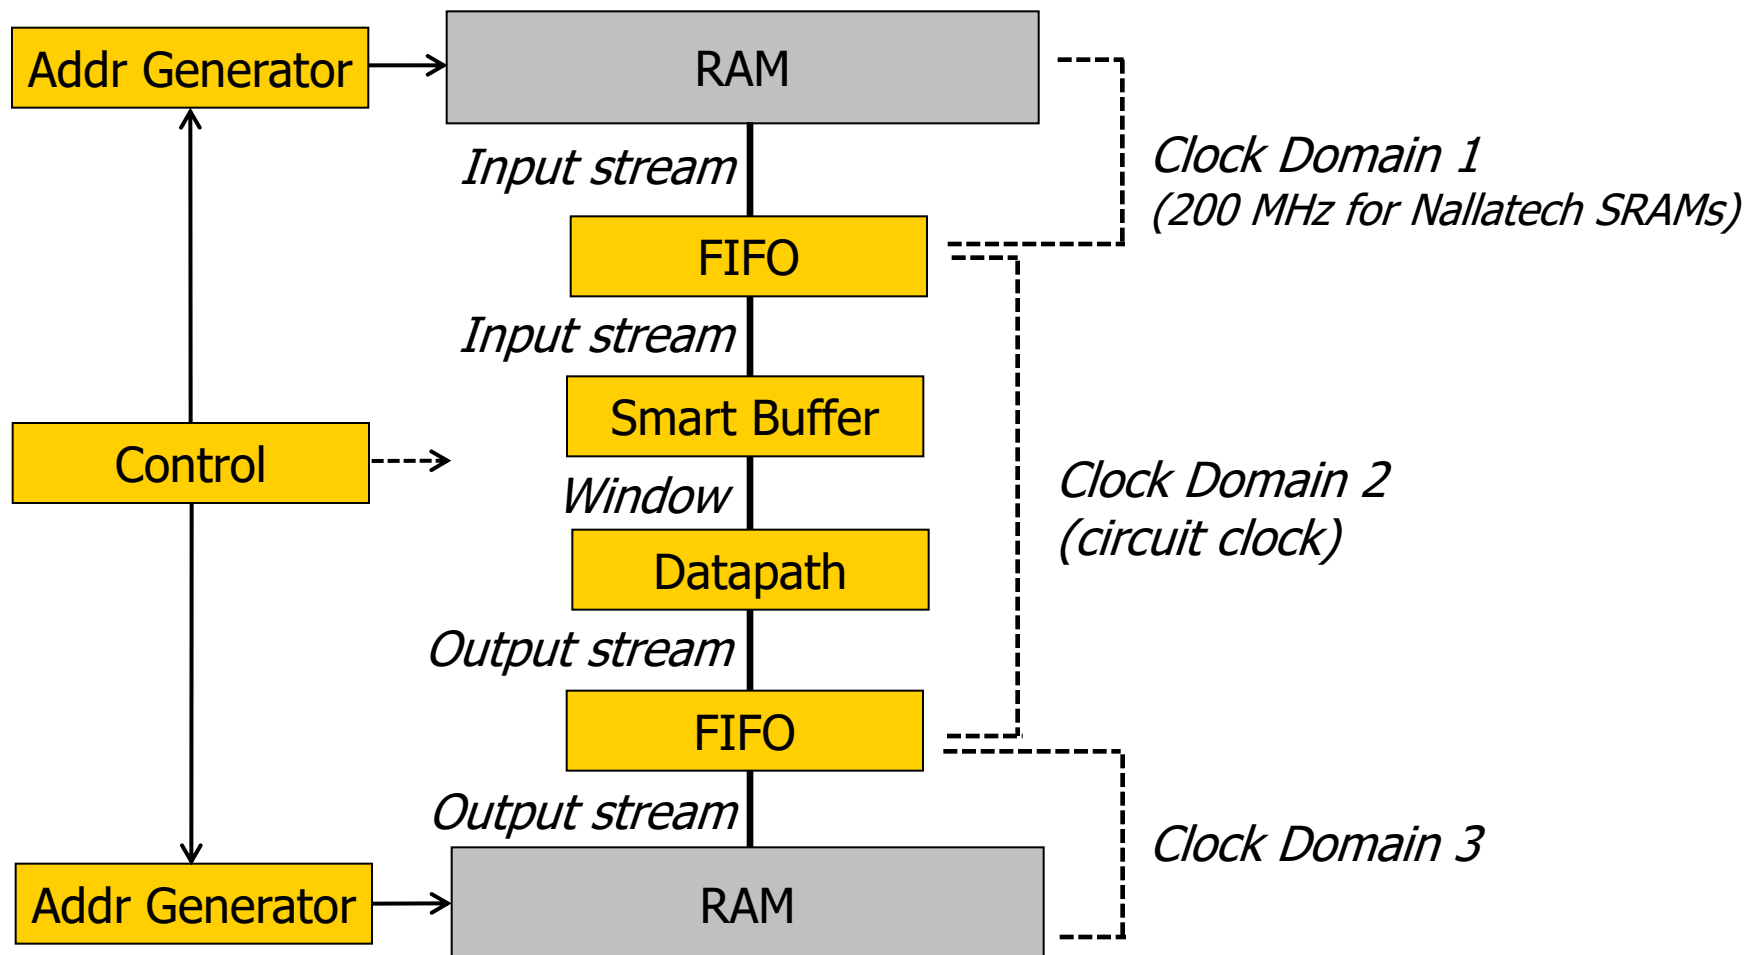
---

- Your turn

```
short b[1004], a[1000];  
for (i=0; i < 1000; i++)  
    a[i] = avg( b[i], b[i+1], b[i+2], b[i+3], b[i+4] );
```

- Analyze memory access patterns
  - Determine window overlap
- Determine smart buffer bandwidth
  - Assume memory bandwidth = 128 bits/cycle
- Determine maximum unrolling with and without smart buffer
- Determine total cycles with and without smart buffer
  - Use previous systolic array analysis (latency, bandwidth, etc).

# Complete Architecture





# Smart Buffer Implementation


---

- Several possibilities (see papers on website)
  - 1) Register smart buffer
  - 2) BRAM smart buffer
- Recommended methodology
  - 1) Determine unrolling amount, which defines window size
  - 2) Determine buffer size and allocate resources
  - 3) Add steering logic to update buffer for each consecutive window

# Register Smart Buffers

```
for (i=0; i < 100; i++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

- 1) Determine unrolling amount
  - Assume FIFO delivers 4 elements/cycle
  - 2 elements reused in each iteration
  - Smart buffer b.w. =  $4+2 = 6$  elements/cycles
- Can execute 4 iterations in parallel
  - Window size = 6 elements
  - 1<sup>st</sup> window = b[0-5], 2<sup>nd</sup> window = b[4-9], etc.



- Important: remember that unrolling also limited by output bandwidth
  - Assume output FIFO can write 4 elements/cycles



# Register Smart Buffers

```
for (i=0; i < 100; i++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

- 2) Determine buffer size and allocate resources
  - Buffer size = # of elements read from FIFO to get first window
  - Assume FIFO provides 4 elements
  - Reading one window (6 elements) requires 2 FIFO accesses
    - 8 total elements read from FIFO
    - Therefore, buffer consists of 8 registers

*Smart buffer:*

Reg

Reg

Reg

Reg

Reg

Reg


Reg

Reg


# Register Smart Buffers

```
for (i=0; i < 100; i++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

- 3) Add steering logic
  - Determine by analyzing access patterns



- Need to shift left by 4




# Register Smart Buffers

```
for (i=0; i < 100; i++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

- Steering logic, cont.
  - But, how does initial window get to the appropriate location?

*1<sup>st</sup> read from FIFO*

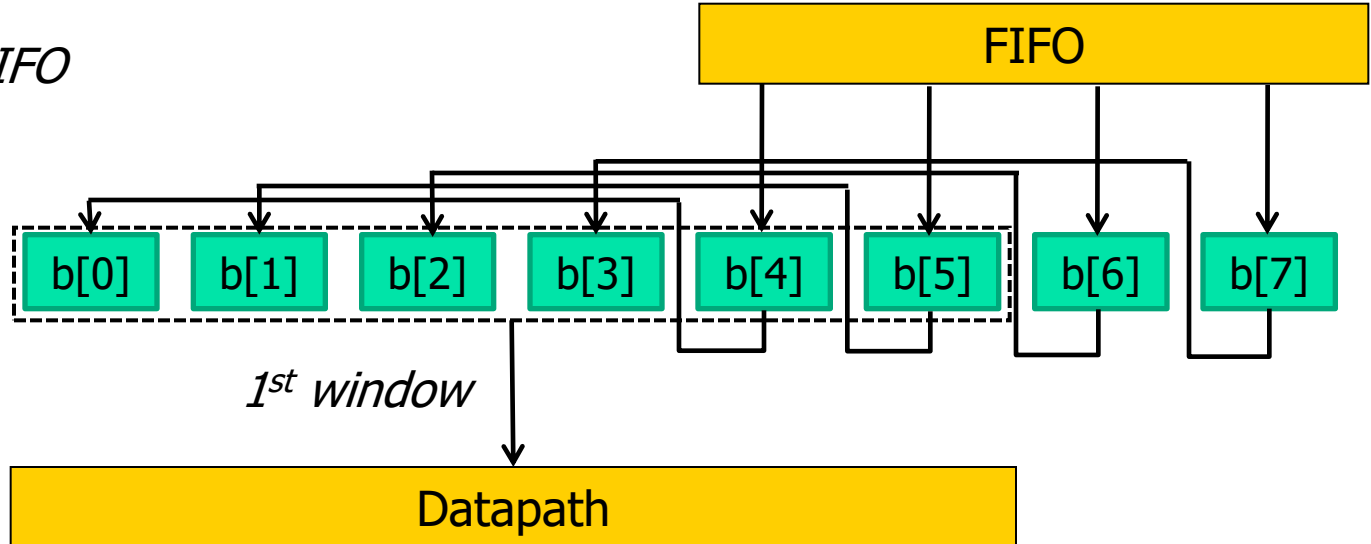


# Register Smart Buffers

```
for (i=0; i < 100; i++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```

- Steering logic, cont.
  - But, how does initial window get to the appropriate location?

*2<sup>nd</sup> read from FIFO*






# Register Smart Buffers

```
for (i=0; i < 100; i++)  
  a[i] = b[i] + b[i+1] + b[i+2];
```


- Steering logic, cont.
  - But, how does initial window get to the appropriate location?

*3<sup>rd</sup> read from FIFO*




# Block RAM Smart Buffers

- Register smart buffers always a possibility
  - But, may have a huge area overhead
- Remember:
  - Buffer size = # of elements read from memory to get first window
- For 1-D examples, usually not much overhead
- Not the case for 2-D examples
- Example:
  - For a 3x3 window, buffer must read first 2 rows of input
  - For an image 1024 pixels wide, smart buffer would require more than 2024 registers




• • • • •



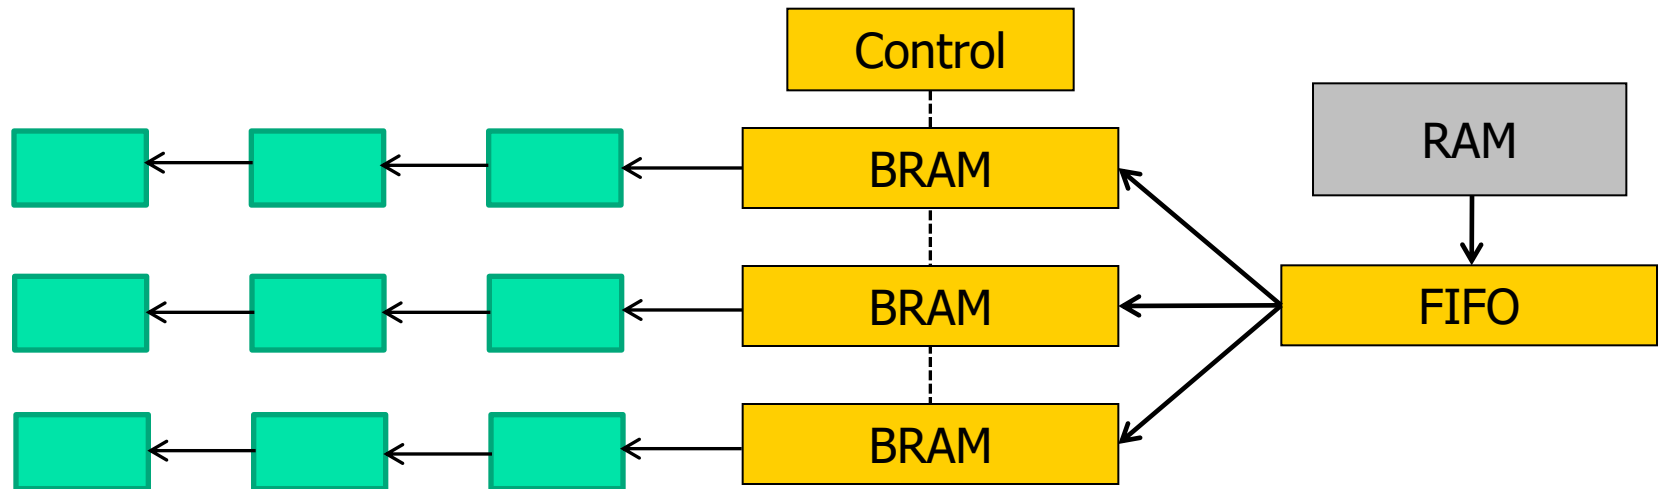
*1<sup>st</sup> 2 rows of input must  
be read before 1<sup>st</sup>  
window is available*

# Block RAM Smart Buffers

- Instead, store each row in block RAM
  - # BRAMs = # of rows in window
  - Size of BRAM = # of columns
- Shift data into a 2-D register array equal to the size of the window
- Example for 3x3 window:

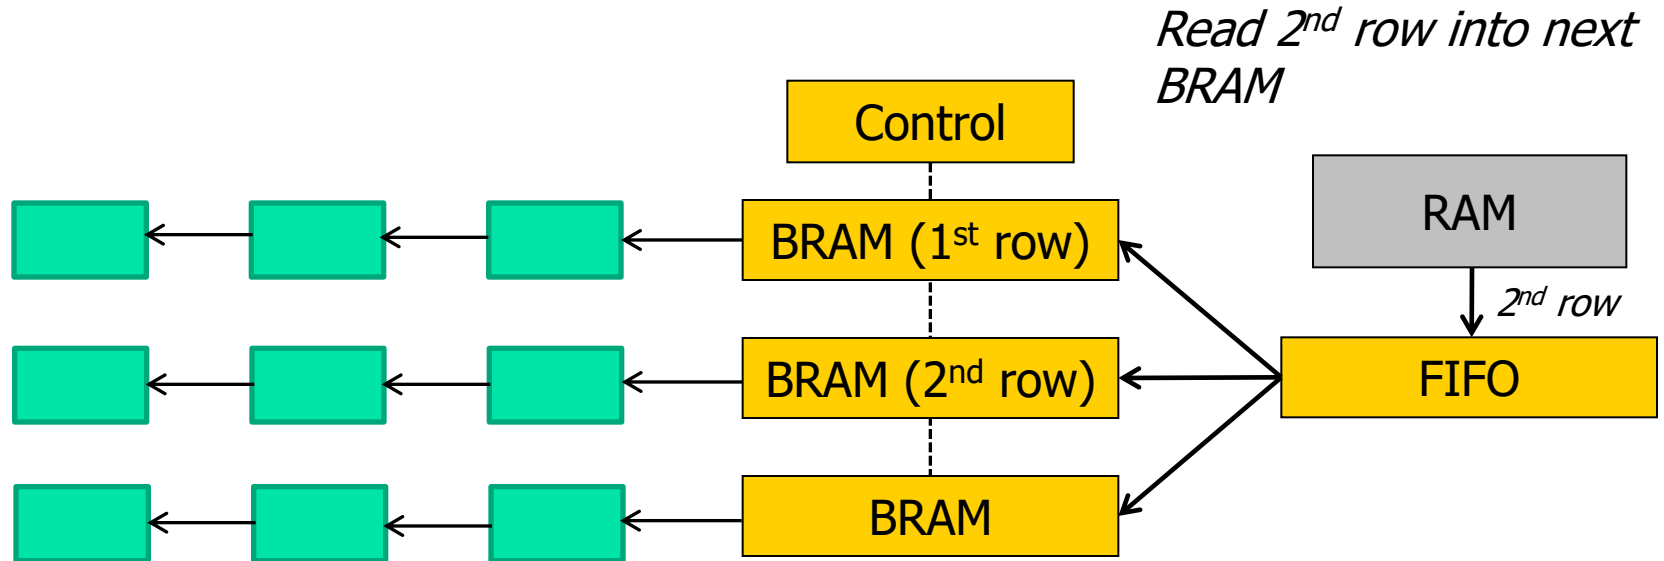


# Block RAM Smart Buffers

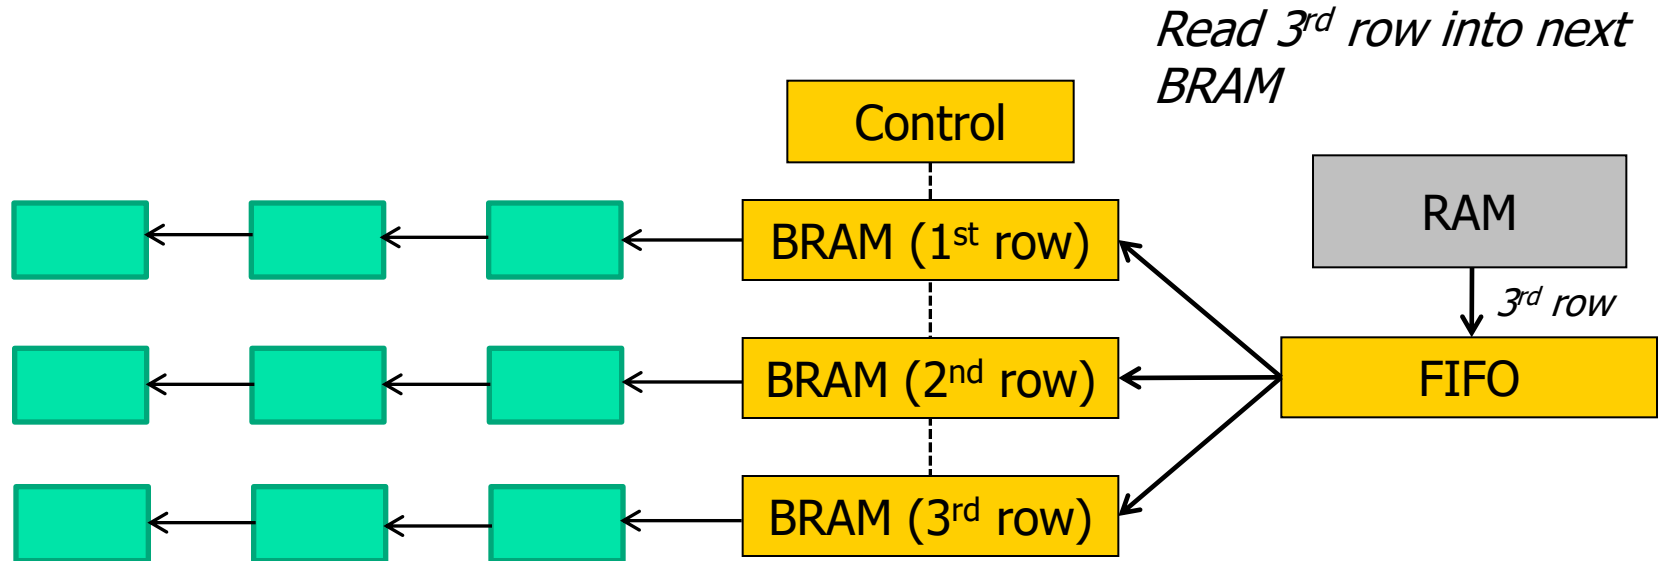




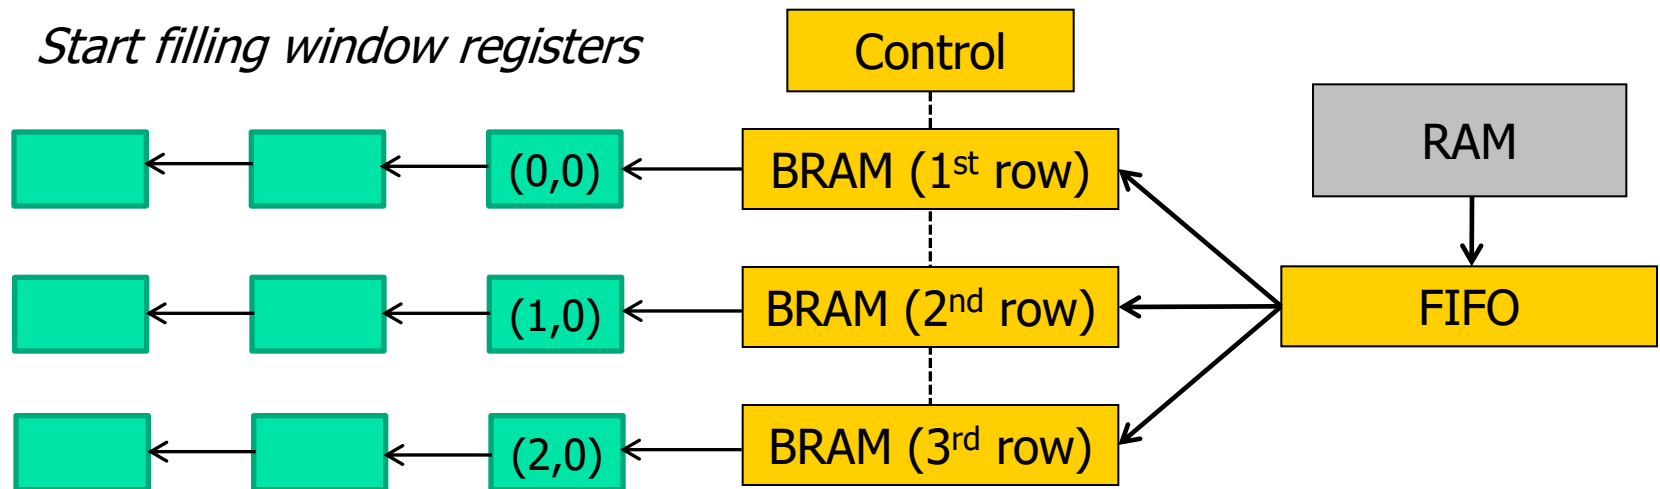
# Block RAM Smart Buffers



# Block RAM Smart Buffers

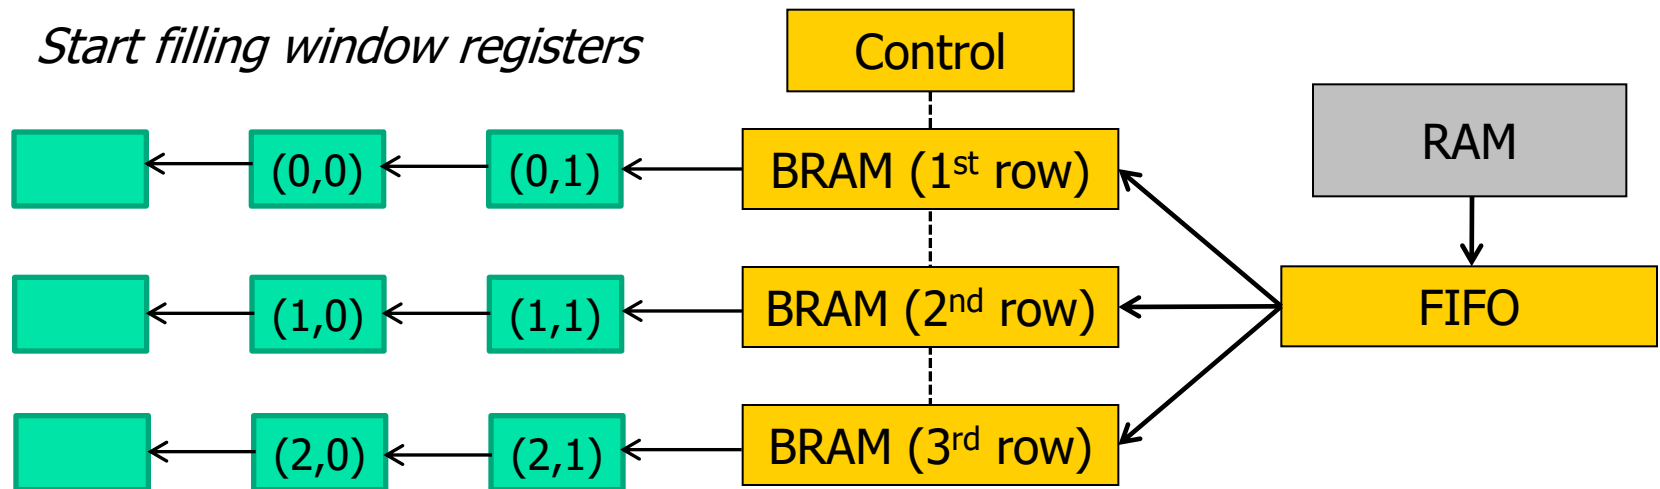


# Block RAM Smart Buffers

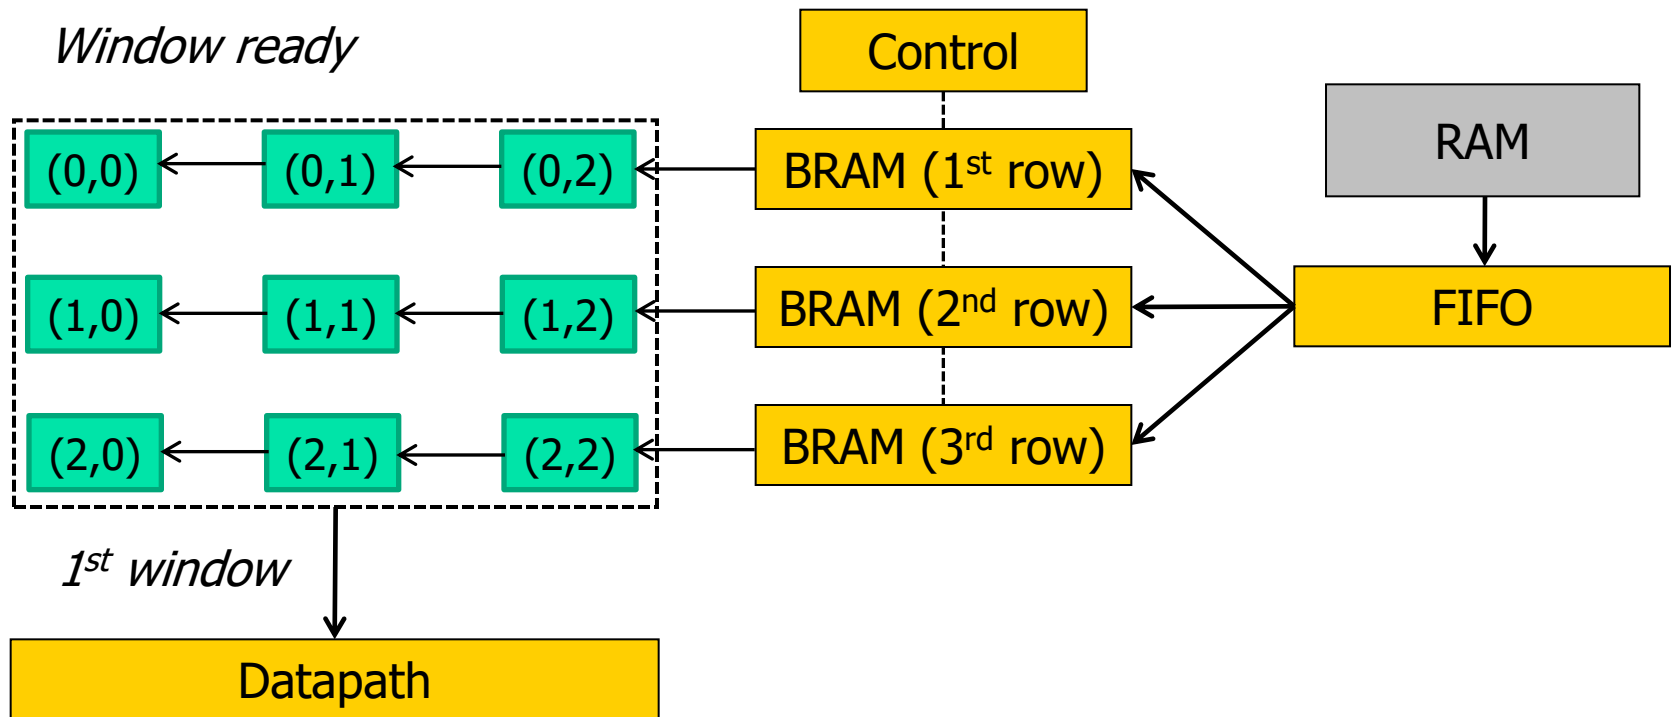




# Block RAM Smart Buffers

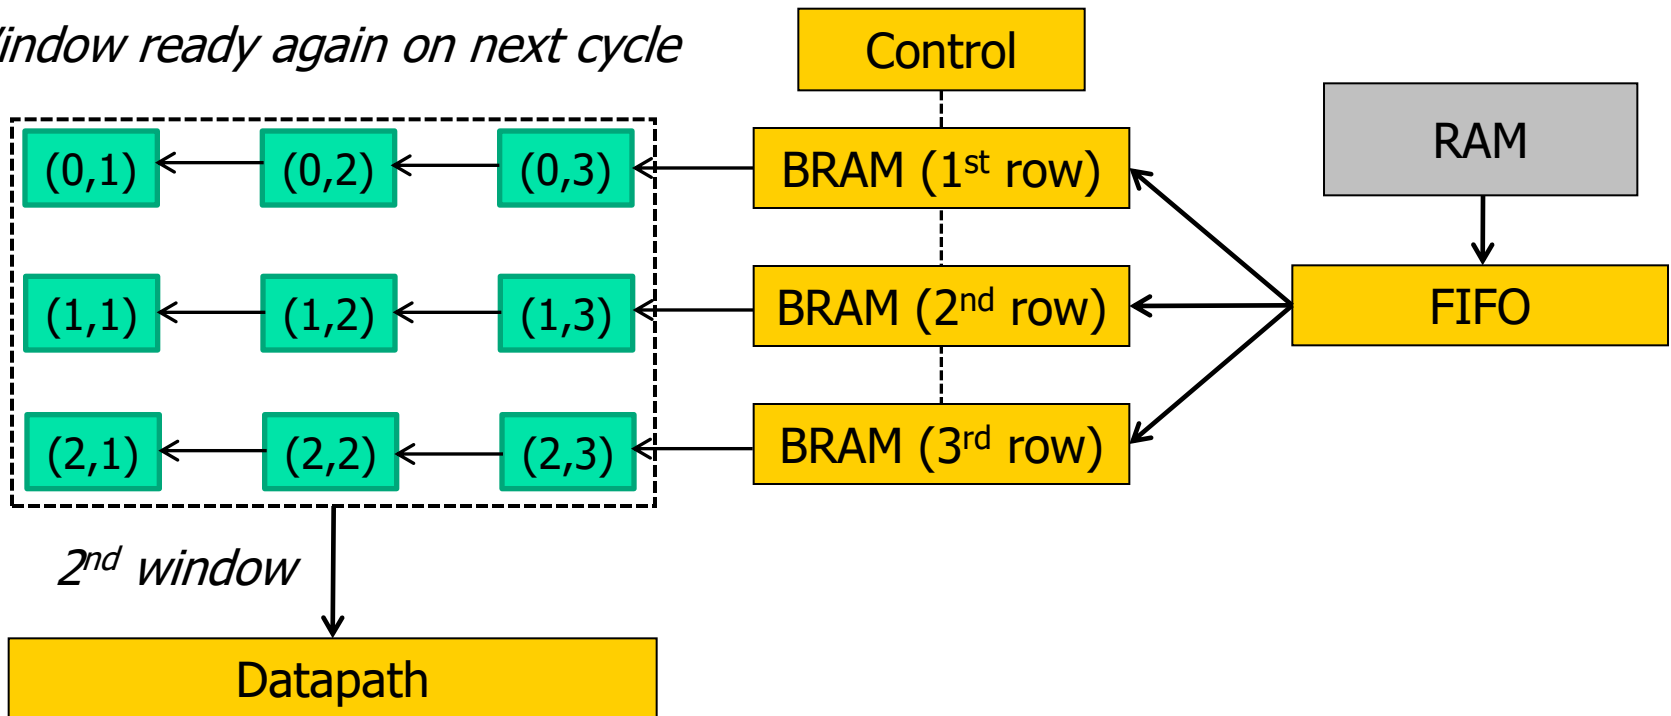


# Block RAM Smart Buffers



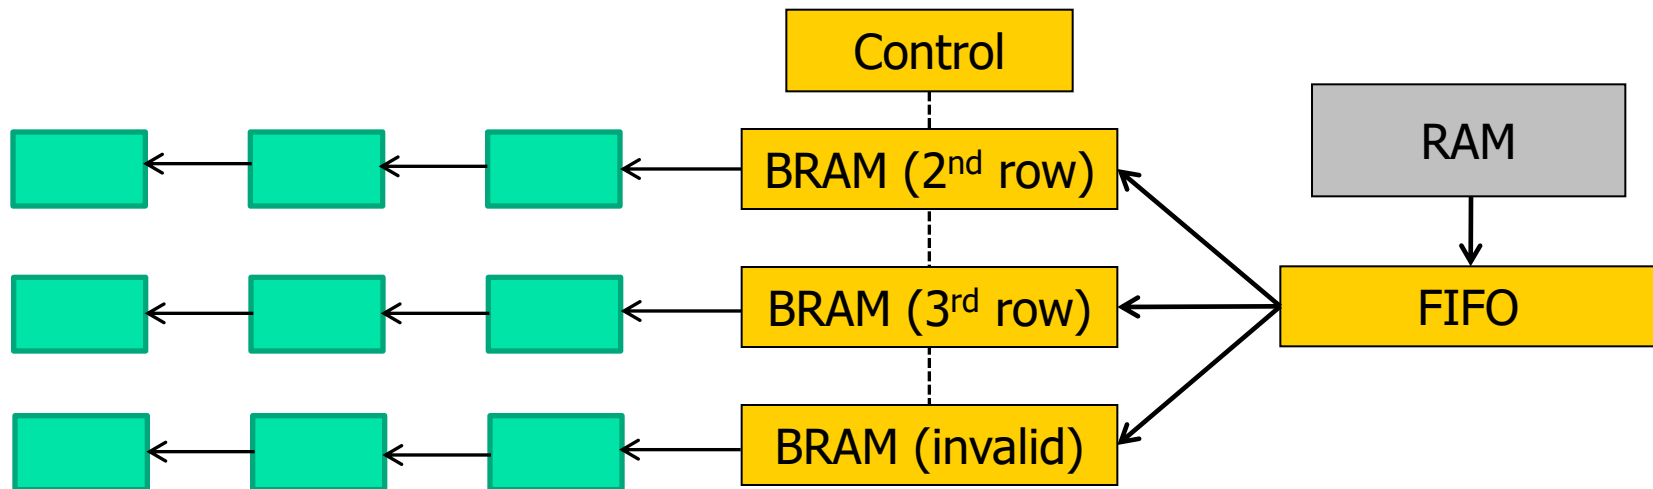
# Block RAM Smart Buffers

*Window ready again on next cycle*



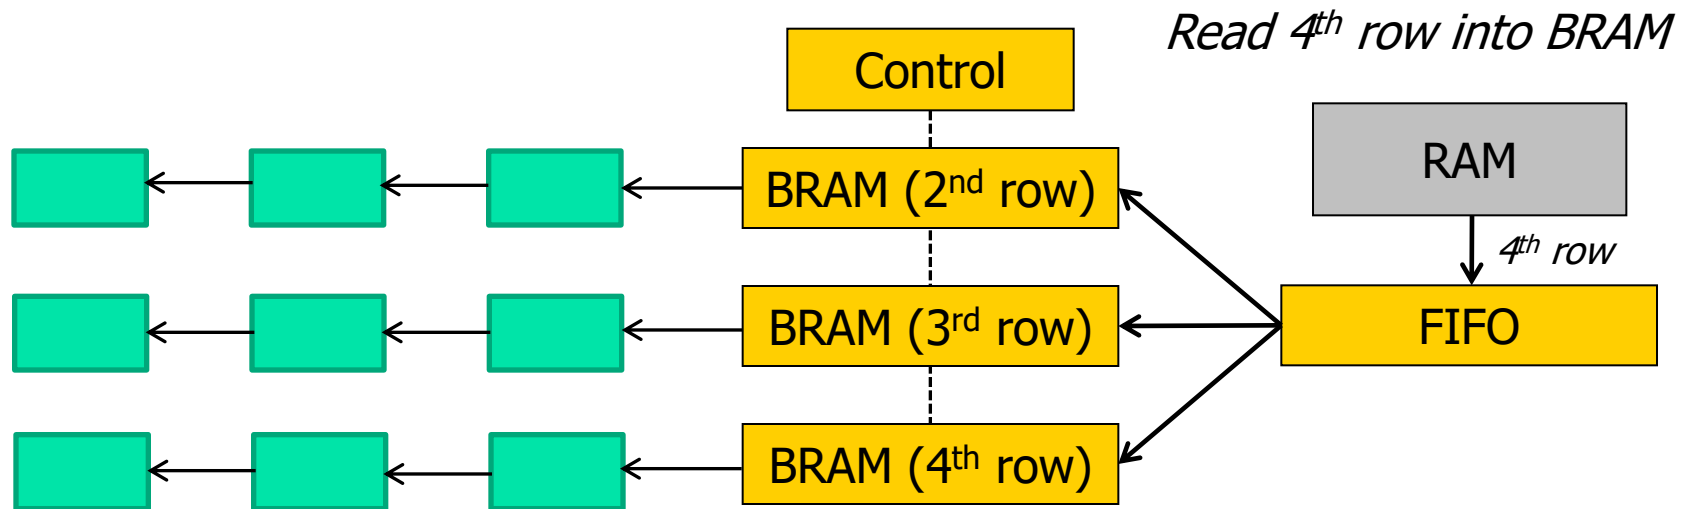
# Block RAM Smart Buffers

*After finishing one row, realign  
BRAMs with next row of windows*

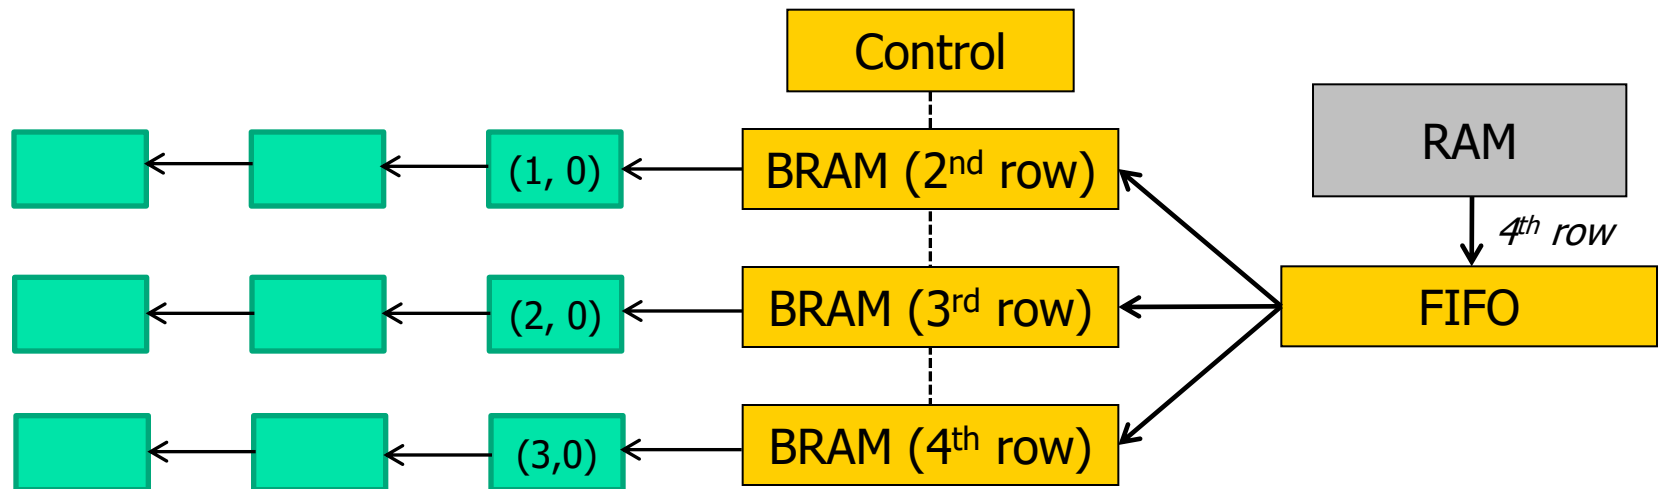


*NOTE: This BRAM alignment is not done by copying data. Instead, the BRAMs are treated as a circular queue, where the front changes after each row*

# Block RAM Smart Buffers



# Block RAM Smart Buffers



*Repeat until entire input has been processed*



# Summary

---

- Smart buffers assemble an input stream into windows
  - Avoids memory accesses for reused data
  - Improves memory bandwidth, enables more unrolling
- Some final thoughts:
  - Smart buffers add latency, but greatly improve throughput
  - For 2-D smart buffers, couldn't you avoid buffering rows by reading elements from the same column of different rows?
    - Yes, but memory bandwidth would be terrible
    - Remember, many memories are efficient for sequential accesses